

Lab 2 — second-site Supervisor deployment

Instantiates the root module for the second lab: **different vCenter, isolated network**, with the subnets swapped relative to lab 1:

Network	Port group	CIDR
Workload	VM Network	192.168.1.0/24
Management	outer-mgmt-net	192.168.2.0/24

Nested ESXi hosts already running at `192.168.1.241-243`. This example has its **own Terraform state** — nothing here touches lab 1 (`examples/lab`), and generated artifacts (TLS cert/key, enable spec) land in `examples/lab2/generated/`, separate from lab 1's.

1. Values you MUST update (search `CHANGE-ME` in `main.tf`)

Variable	What to set it to
<code>vcenter_server</code>	The lab-2 vCSA's system name (PNID) — FQDN or IP , whichever it was deployed with (the CN of its TLS cert tells you). If an FQDN: it must resolve to the vCSA's IP from this machine — verify with <code>dscacheutil -q host -a name <fqdn></code> (macOS), not just <code>dig</code> . If an IP: use it as-is; this also sidesteps DNS-hijack risk entirely
<code>datacenter</code>	Datacenter name in the lab-2 vCenter (default <code>Datacenter</code>)
<code>physical_host_name</code>	The physical host's inventory name in vCenter — typically its IP
<code>physical_host_cluster</code>	Cluster containing the physical host (default <code>Cluster</code>)
<code>supervisor_cluster</code>	Cluster containing the three nested hosts (default <code>Supervisor-Cluster</code>)
<code>outer_datastore</code>	Datastore on the physical host where the HAProxy + NFS VMs deploy (default <code>datastore1</code>)

Variable	What to set it to
<code>management_gateway</code> / <code>management_dns</code>	Router + DNS reachable from <code>192.168.2.0/24</code> at this site
<code>workload_gateway</code> / <code>workload_dns</code>	Router + DNS reachable from <code>192.168.1.0/24</code> at this site

Verify-but-probably-keep (defaults follow the lab-1 conventions):

Variable	Default	Check
<code>nested_esxi_hosts</code>	<code>192.168.1.241-243</code>	must equal the hosts' vCenter inventory names exactly
<code>management_cp_starting_ip</code>	<code>192.168.2.231</code>	5 consecutive IPs free?
<code>nested_host_mgmt_ips</code> (in the module block)	<code>.1.24x → .2.24x</code>	mgmt IPs free? keys match <code>nested_esxi_hosts</code> ?
<code>workload_ip_range</code>	<code>192.168.1.201-230</code>	outside DHCP, doesn't overlap the router/hosts/VMs
<code>vip_pool</code> / <code>vip_pool_usable</code>	<code>192.168.1.248/29 → .249-.254</code>	free, outside DHCP
<code>nfs_ip</code> / <code>haproxy_ip</code> (module block / variable)	<code>.1.244</code> / <code>.1.245</code>	free
<code>ntp_servers</code> (module block)	<code>162.159.200.1</code>	reachable on UDP 123 from the physical host

Secrets go in `secrets.auto.tfvars` (never committed — see step 3).

2. Prerequisites before `terraform apply`

On the nested ESXi hosts (do this **FIRST**):



Set a unique hostname on each host — see §2a below for the exact commands. ESXi defaults to `localhost`; with that, every spherelet receives the same client-cert identity (`system:node:localhost`) and Kubernetes blocks node registration — the

cluster comes up with zero worker nodes. See `TROUBLESHOOTING.md` → “Supervisor ESXi nodes never join”.



Each host has **three vNICs**: vmnic0 + vmnic1 on the outer `VM Network`, vmnic2 on the outer `outer-mgmt-net` path. ESXi does not detect hot-added NICs — power-cycle after adding.



Expected host networking at this stage (naming trap): the host’s own “**Management Network**” (`vmk0` on `vSwitch0` / vmnic0) lives on the **workload** CIDR — `192.168.1.24x` — because that’s the host’s address. The `192.168.2.0/24` “management” subnet is management for the *Supervisor control plane*, not the hosts; the host only gets a leg on it when Terraform adds `vmk1` (sup-host-mgmt). Leave `vmk0` where it is, and leave vmnic1/vmnic2 unattached — the DVS claims them.



All three hosts joined to the Supervisor cluster and `connected`.



Cluster vLCM image matches the installed ESXi build (vSphere UI → Cluster → Updates → Image → “All hosts compliant”). Mismatch fails Supervisor enable with VIB download errors.

In the lab-2 vCenter:



Datacenter, both clusters, physical host, and datastore exist under the names configured above. **Terraform creates none of these** — every `datacenter` / `*_cluster` / `*_host` / `*_datastore` variable is a lookup of existing inventory, and `plan` fails if a name doesn’t match. In particular, create the Supervisor cluster and join the three nested hosts to it by hand (vSphere Client → New Cluster → drag hosts in, or `govc cluster.create` + `govc cluster.add`).



Port groups `VM Network` and `outer-mgmt-net` exist on the physical host (Terraform sets their security flags — you don’t need to).



SSO admin credentials for the account in `secrets.auto.tfvars`.

Network / environment:



Routing between `192.168.1.0/24` and `192.168.2.0/24` (the Supervisor enable validates cross-subnet DNS from the CP VM).



All static IPs in §1 are outside any DHCP scope.



WAN egress from the workload subnet — the Ubuntu cloud OVA and the HAProxy Dataplane API binary download at deploy time.



`vcenter_server` resolves correctly from this machine **and** will resolve correctly from the CP VMs (via `management_dns`). A public CNAME on the vCenter name will poison both — see `TROUBLESHOOTING.md` → DNS resolution.



Clocks sane on the physical host and vCSA (Terraform enables NTP on the physical host, but a wildly wrong vCSA clock still breaks TLS).

On this machine (once per workstation):



`make -C ../../ install-deps` — terraform, govc, python3+pyvmomi, curl, jq, openssl.

2a. Setting the ESXi hostnames (required — do before Terraform)

What to set: the recommended values for this lab are

Host	<code>--host</code>	<code>--domain</code>
192.168.1.241	<code>nested-esxi-1</code>	<code>lab2.skynetsystems.io</code>
192.168.1.242	<code>nested-esxi-2</code>	<code>lab2.skynetsystems.io</code>
192.168.1.243	<code>nested-esxi-3</code>	<code>lab2.skynetsystems.io</code>

The actual strings are your choice — the rules that matter:

- **Unique per host** (the whole point — `localhost` on all three is the failure mode).
- **Lowercase letters, digits, hyphens, dots only** — the FQDN becomes the Kubernetes node name, which must be DNS-1123 compliant (no underscores, no uppercase).
- **Stable** — spherelet's cert identity is minted from it; renaming later means re-issuing spherelet certs.
- **Avoid** `.local` as the domain (mDNS conflicts, especially from macOS).
- No DNS records are required for these names; the domain is effectively a label.

Pick whichever method fits:

Helper script (from the repo root; checks first, sets only when wrong, safe to re-run — needs SSH enabled on the hosts, see the SSH primer in §3):

```
cd ~/Repos/vcf9-supervisor-terraform
SSHPASS='<esxi root password>' ./scripts/sv-set-hostnames -d <lab2-domain> \
  192.168.1.241=nested-esxi-1 \
  192.168.1.242=nested-esxi-2 \
  192.168.1.243=nested-esxi-3
```

Manual over SSH, one host at a time:

```
ssh root@192.168.1.241 'esxcli system hostname set --host=nested-esxi-1 --
  domain=<lab2-domain>'
ssh root@192.168.1.241 'esxcli system hostname get' # verify
```

ESXi Host Client (no SSH needed): browse to `https://<host-ip>/ui`, log in as root → **Networking** → **TCP/IP stacks** → **Default TCP/IP stack** → **Edit settings** → set Host name + Domain.

DCUI (VM console): F2 → log in → **Configure Management Network** → **DNS Configuration** → set Hostname.

Setting the hostname does **not** rename the host in vCenter inventory (that stays the IP it was added with) and does not disrupt the host.

If spherelet ever ran on a host while it was named `localhost`, also clear `/etc/vmware/spherelet/*.crt|*.key|*.pem` on it and restart spherelet after Supervisor enable (full recipe in TROUBLESHOOTING.md).

2b. Creating the Supervisor cluster (manual — Terraform won't do this)

Settings below are the lab-1 working configuration, read live from its cluster.

Via the vSphere Client

1. **Create the cluster**: right-click the datacenter → **New Cluster** → name it `Supervisor-Cluster` (or whatever you set `supervisor_cluster` to), with:

Setting	Value	Why
vSphere DRS	On , automation level Fully Automated	Supervisor requires DRS to place CP/pod VMs

Setting	Value	Why
vSphere HA	On (admission control on)	Supervisor requires HA
vSAN	Off	we use the NFS datastore instead

(Don't look for a "Manage all hosts with a single image" checkbox — that's the vSphere 7/8 wizard. In vSphere 9, image-based lifecycle management is the only mode, so there's nothing to tick; the image itself is configured after creation, next step.)

2. **Set the cluster image to the hosts' exact ESXi build** — do this after creating the cluster and BEFORE enabling Supervisor. Enable pushes the cluster image (with the spherelet VIBs) onto the hosts; any version mismatch fails with `Cannot download VIB` (runbook Root Causes #1–2).
 1. Import the offline depot matching the build the nested hosts actually run (e.g. `VMware-ESXi-9.0.x-<build>-depot.zip` from Broadcom support):
vSphere Client → **Lifecycle Manager** → **Import** → **Bundle (.zip)**.
 2. **Cluster** → **Updates** → **Image** → **Edit** → **ESXi Version** dropdown → pick that exact build → **Validate** → **Save**.
 3. Target state: **"All hosts compliant."** Don't proceed until it is.
3. **Add the hosts:** right-click the cluster → **Add Hosts** → enter `192.168.1.241/.242/.243` with each host's `root` password → accept the thumbprints. (Order note: adding hosts before or after setting the image both work — compliance is evaluated once both exist.)

The wizard leaves the hosts in maintenance mode — this is expected. Exit it on each host (right-click → Maintenance Mode → Exit, or):

```
for h in 192.168.1.241 192.168.1.242 192.168.1.243; do
  govc host.maintenance.exit /Datacenter/host/Supervisor-Cluster/$h
done
```

Within ~60 s of leaving maintenance mode, vSphere auto-deploys two small **vCLS** VMs on the cluster — normal, DRS needs them, leave them alone. HA datastore/network alarms at this stage are also expected until step 4's advanced options are set and Terraform mounts `nfs-shared`.

4. **HA advanced options** (silences recurring HA alarms in a nested, single-datastore lab — set via **Cluster** → **Configure** → **vSphere Availability** → **Edit** → **Advanced Options**):

Option	Value	Why
<code>das.ignoreInsufficientHbDatastore</code>	<code>true</code>	HA wants ≥ 2 heartbeat datastores; nested hosts will only ever see <code>nfs-shared</code>
<code>das.ignoreRedundantNetworkWarning</code>	<code>true</code>	HA wants redundant management NICs; the nested hosts have one path

(See `TROUBLESHOOTING.md` → “HA alarm spam” for a pyvmomi snippet that sets these from the CLI — govc doesn’t expose them.)

Or via govc

```
# Point govc at the lab-2 vCenter:
export GOVC_URL=<lab2-vcenter-fqdn> GOVC_USERNAME=administrator@vsphere.local \
    GOVC_PASSWORD='<ss0 password>' GOVC_INSECURE=true

govc cluster.create -dc=Datacenter Supervisor-Cluster

govc cluster.change -drs-enabled -drs-mode=fullyAutomated -ha-enabled \
    /Datacenter/host/Supervisor-Cluster

for h in 192.168.1.241 192.168.1.242 192.168.1.243; do
    govc cluster.add -cluster=Supervisor-Cluster \
        -hostname=$h -username=root -password='<esxi root pw>' -noverify
done

# If any host was added in maintenance mode:
govc host.maintenance.exit /Datacenter/host/Supervisor-Cluster/192.168.1.241

# Verify: all three connected
govc find /Datacenter/host/Supervisor-Cluster -type h
```

The vLCM image assignment and the two HA advanced options still need the UI / pyvmomi as noted above.

Don’t configure networking on the cluster beyond this — the `supervisor-dvs`, port groups, uplinks, and vmkernel NICs are all Terraform’s job.

3. Commands to run (in order)

All paths relative to the repo root (`~/Repos/vcf9-supervisor-terraform`); the example lives in `examples/lab2/`.

How to SSH into an ESXi host (step 0 needs this)

Account: always the user `root`, with the ESXi root password chosen when the host was installed (this is a per-host password — not a vCenter/SSO login, and not the `secrets.auto.tfvars` passwords).

SSH is disabled by default on ESXi. Enable it first, any one of these ways:

- **vSphere Client:** select the host → Configure → System → Services → SSH → Start.
- **ESXi Host Client:** browse to `https://<host-ip>/ui`, log in as root → Manage → Services → `TSM-SSH` → Start.
- **DCUI** (the yellow console screen on the VM): F2 → log in → Troubleshooting Options → Enable SSH.
- **govc** (from this machine, pointed at the lab-2 vCenter): `GOVC_HOST=/
<datacenter>/host/<cluster>/<host-ip> govc host.service start TSM-SSH`

Running a command over SSH — either interactively:

```
ssh root@192.168.1.241      # type the root password when prompted
esxcli system hostname get  # you're now in the ESXi shell
exit
```

or one-shot without opening a shell (this is the form the commands below use — everything inside the quotes runs on the host):

```
ssh root@192.168.1.241 'esxcli system hostname get'
```

When done, stop SSH again from the same menu (or `govc host.service stop TSM-SSH`) — leaving it on triggers a persistent yellow warning on the host and is poor hygiene.

```
# 0. FIRST: verify each nested host has a unique hostname — ESXi
#   installs default to "localhost", which silently breaks spherelet
#   node registration (§2 / TROUBLESHOOTING.md). The helper script
#   checks and only sets when needed (idempotent, safe to re-run):
cd ~/Repos/vcf9-supervisor-terraform
SSHPASS='<esxi root password>' ./scripts/sv-set-hostnames -d <lab2-domain> \
  192.168.1.241=nested-esxi-1 \
  192.168.1.242=nested-esxi-2 \
  192.168.1.243=nested-esxi-3
#   (Omit SSHPASS to be prompted per host. If SSH is disabled on a
#   host, enable it via govc against the lab-2 vCenter first:
#   GOVC_HOST=/<datacenter>/host/<supervisor-cluster>/<host-ip> \
```

```
#      govc host.service start TSM-SSH
#      — or use the DCUI/host client. Equivalent manual command:
#      ssh root@<host> 'esxcli system hostname set --host=<name> --
      domain=<domain>')

cd ~/Repos/vcf9-supervisor-terraform/examples/lab2

# 1. Secrets file (gitignored) — copy the template and fill in:
cp secrets.auto.tfvars.example secrets.auto.tfvars
vi secrets.auto.tfvars          # vcenter_password, haproxy_password, vcenter_ip
chmod 600 secrets.auto.tfvars

# 2. Fill in the CHANGE-ME values:
vi main.tf

# 3. Initialize providers (first run only):
terraform init

# 4. Sanity-check the config:
terraform validate

# 5. Preview — expect ~20+ resources to create, 0 to destroy:
terraform plan

# 6. Apply (~25-30 min unattended; the Supervisor-enable step polls
#    up to 45 min at the end):
terraform apply

# 7. Post-apply outputs (kubectl login instructions, API VIP):
terraform output next_steps
```

Or drive it through the Makefile from the repo root:

```
cd ~/Repos/vcf9-supervisor-terraform
make TF_DIR=examples/lab2 init
make TF_DIR=examples/lab2 plan
make TF_DIR=examples/lab2 apply
```

(Note: `make hard-check` and the `scripts/sv-*` helpers are hardwired to lab 1's IPs/vCenter — don't use them against this site without editing.)

Teardown: `terraform destroy` from `examples/lab2/` — disables the Supervisor cleanly first (10–15 min), then removes the VMs, DVS, and policy.

4. After apply — verify

```
# All 4-6 nodes Ready (3 CP + your 3 ESXi agents). CP root password:
# ssh to the lab-2 vCSA → shell → /usr/lib/vmware-wcp/decryptK8Pwd.py
ssh root@<cp-floating-ip> 'kubectl get nodes'
```

```
# API reachable through the HAProxy VIP:  
curl -sk --max-time 5 "${cd examples/lab2 && terraform output -raw  
    supervisor_api_vip)/version}"
```

If the ESXi agent nodes are missing, work through [TROUBLESHOOTING.md](#) → “Supervisor ESXi nodes never join” — check the spherelet cert identity (hostname issue) first, then the vmk path.